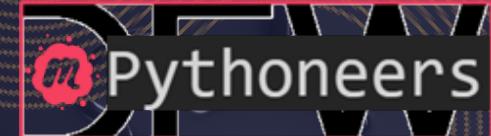




Real-Time, Data-Driven Decision-Making with Spark Streaming

May 4, 2023





AGENDA

01 From Big Data to Databricks

02 Streaming Analytics Basics

03 Demo, Demo, Demo!!!

SPEAKER INFO

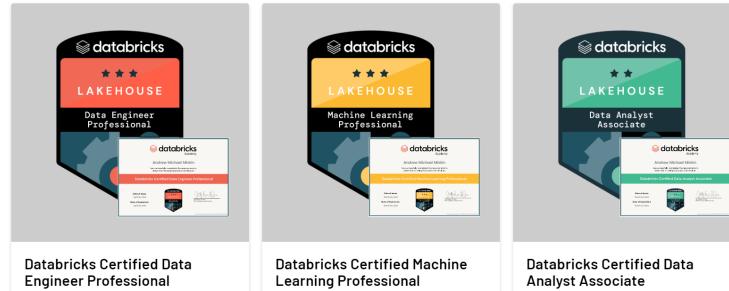
Previous Incarnations

- SQL Support/Consulting 2000-06
- Data Scientist 2010 – Now
- 5 Startups 2007-19
- Analytics Architect/Dev
- Entrepreneur



Current Incarnation

- Advisory Data/AI Architect





From Big Data to Databricks

Why Big Data?



THIN THE LINES
BETWEEN
OPERATIONAL AND
ANALYTICAL

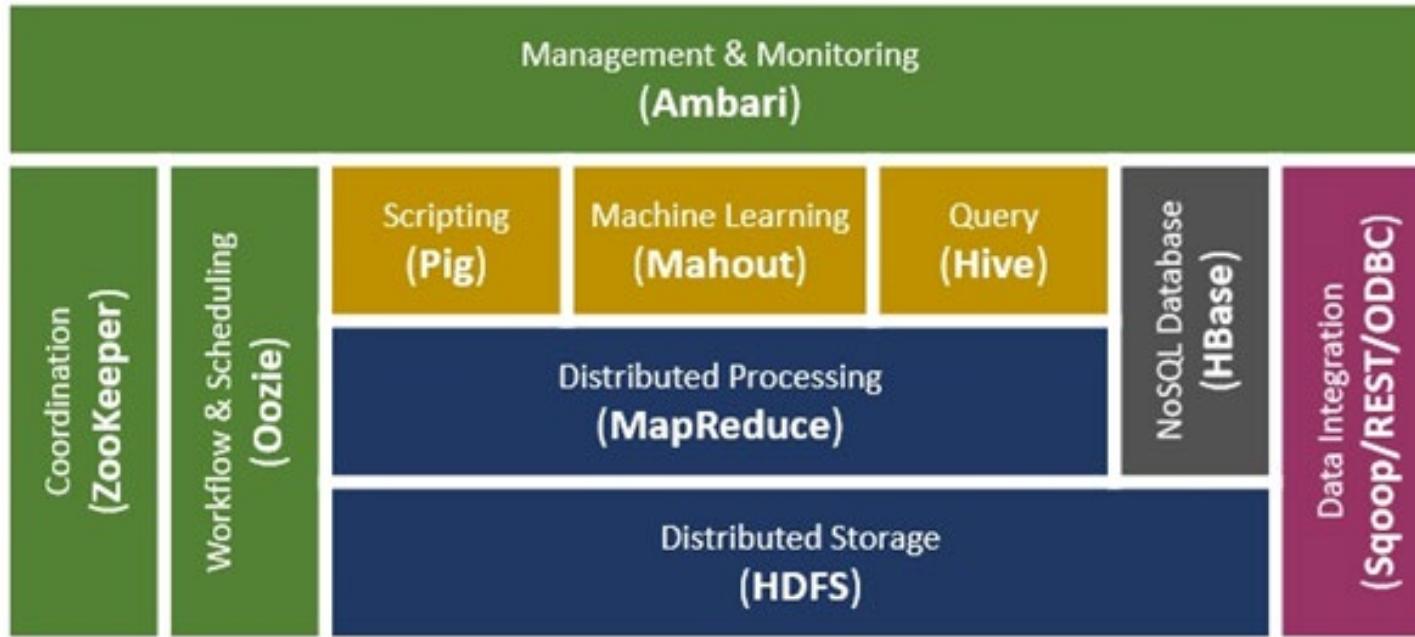


DECENTRALIZED
WORKSPACES
CENTRALIZED STORAGE

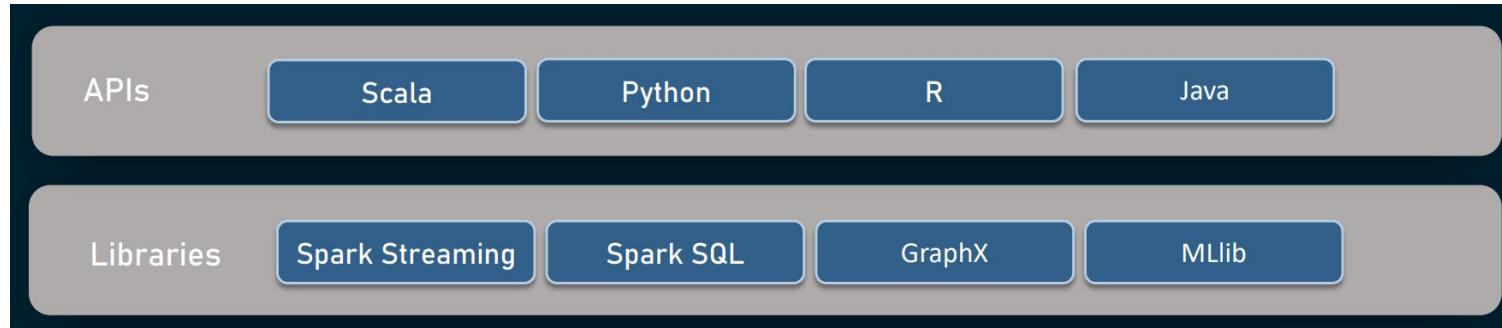


COMPUTE IS
PAY AS YOU GO

Hadoop Architecture (Pre-Spark)



Spark APIs



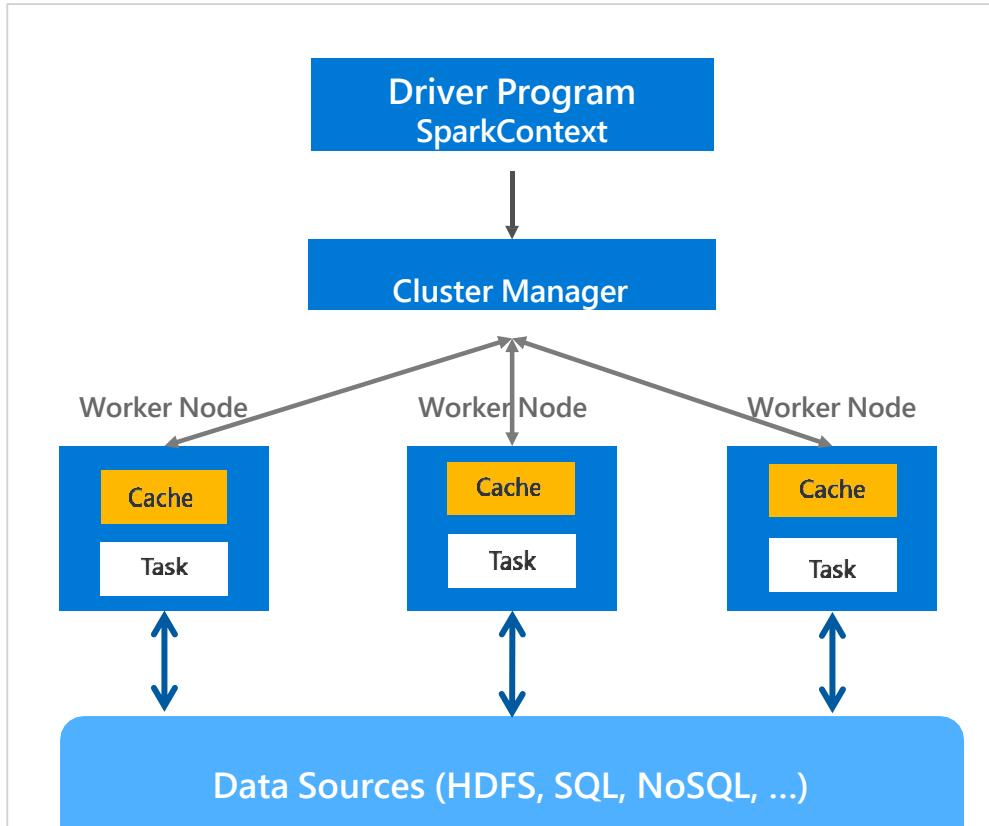
SEPARATION OF COMPUTE AND STORAGE

NATIVE JVM WITH SCALA DIALECT

SPARKR AND PYSPARK THUNKY INTEGRATIONS

Spark Architecture

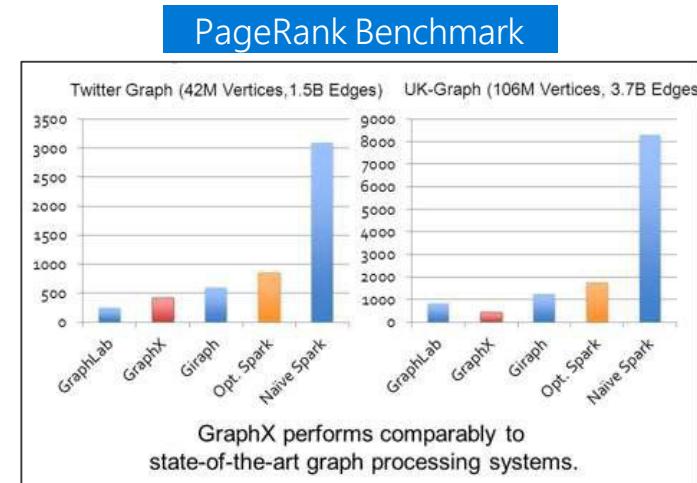
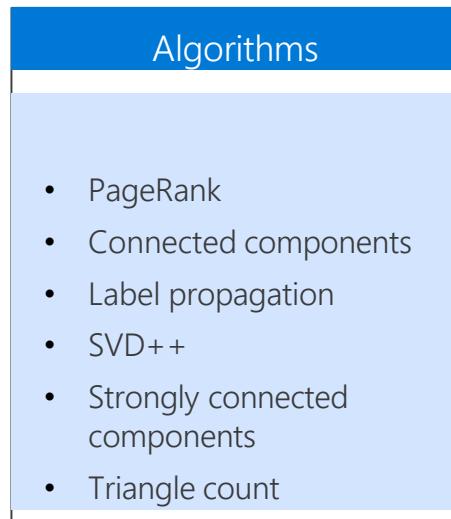
- 'Driver' runs the user's 'main' function and executes the various parallel operations on the worker nodes.
- The results of the operations are collected by the driver
- The worker nodes read and write data from/to Data Sources including HDFS.
- Worker node also cache transformed data in memory as RDDs (Resilient Data Sets).
- Worker nodes and the Driver Node execute as VMs in public clouds (AWS, Google and Azure).



GraphX

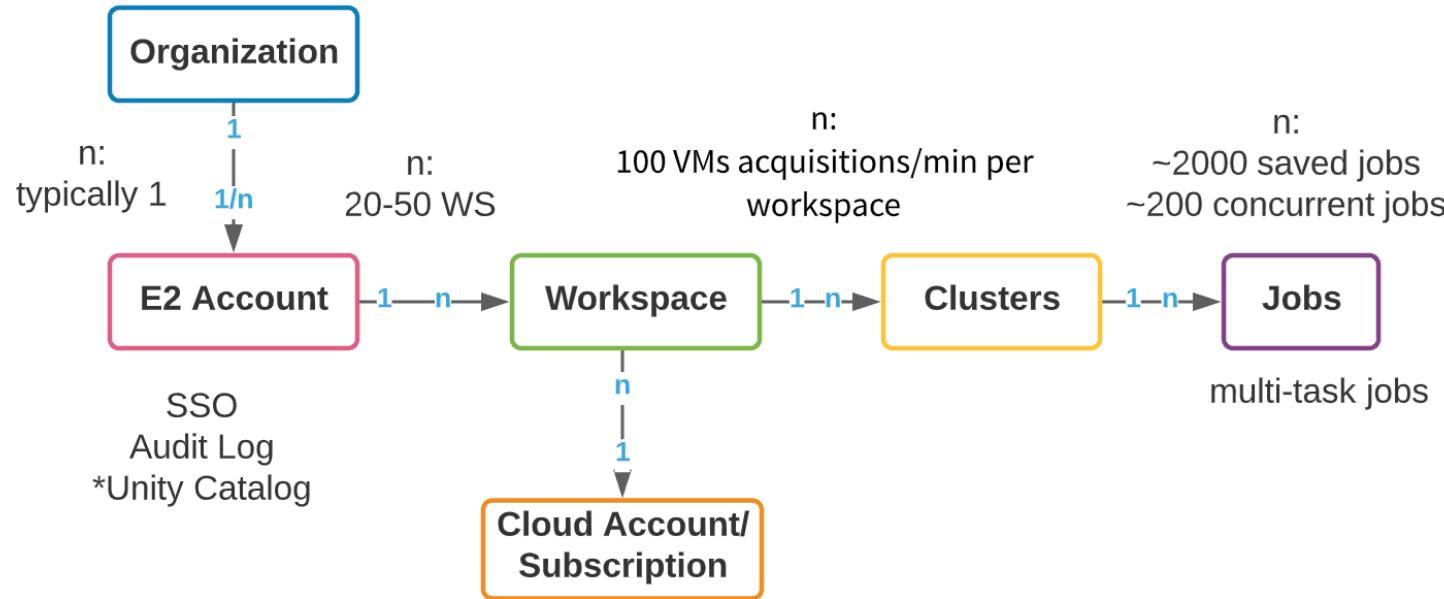
A set of APIs for graph and graph-parallel computation.

- Unifies ETL, exploratory analysis, and iterative graph computation within a single system.
- Developers can:
 - [view](#) the same data as both graphs and collections,
 - [transform](#) and [join](#) graphs with RDDs, and
 - write custom iterative graph algorithms using the [Pregel API](#).
- Currently only supports using the Scala and RDD APIs.



Source: [AMPLab](#)

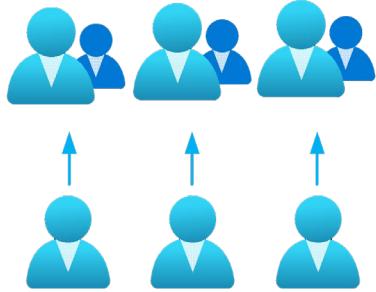
Databricks High Level Object Model



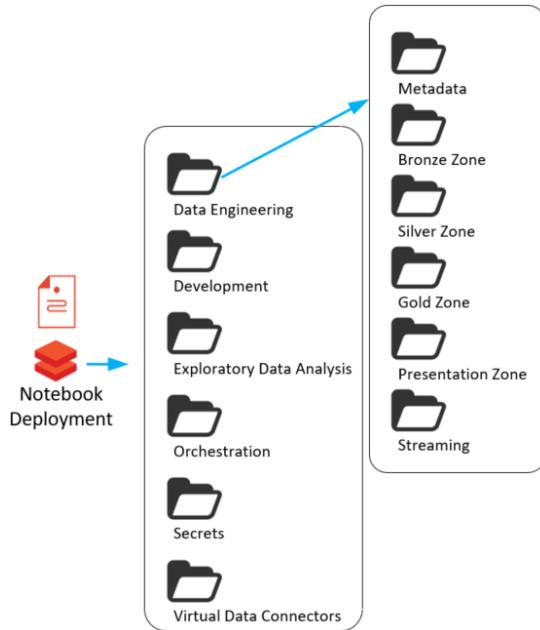
1 cloud account can be used for multiple workspace

Databricks Configuration

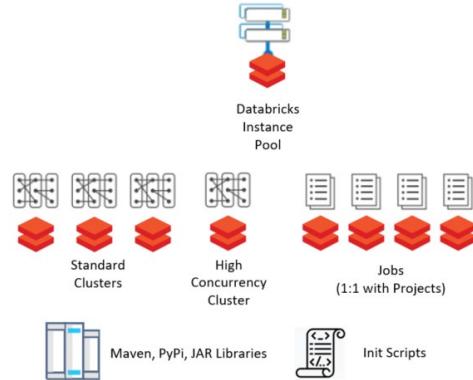
Databricks Security & IAM Roles



Databricks Workspace/Notebooks



Pools, Clusters, Jobs, Libraries, Init Scripts

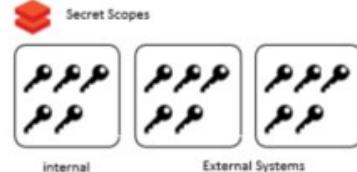


Upload Prerequisites



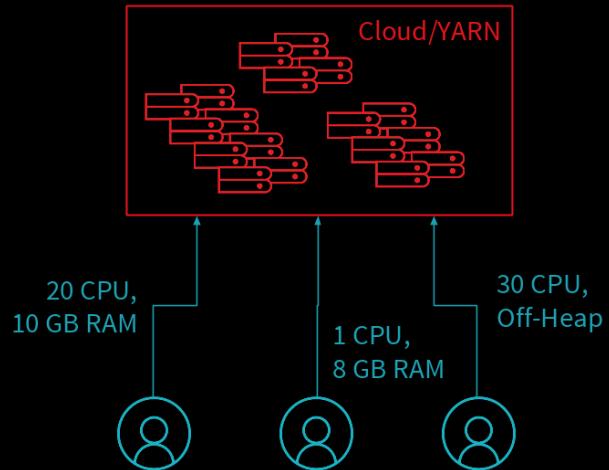
DBFS
ADLS
01100
10010
00101
Seed Data
Static Data
JARs

Secrets



Databricks Serverless

Current Spark Environments



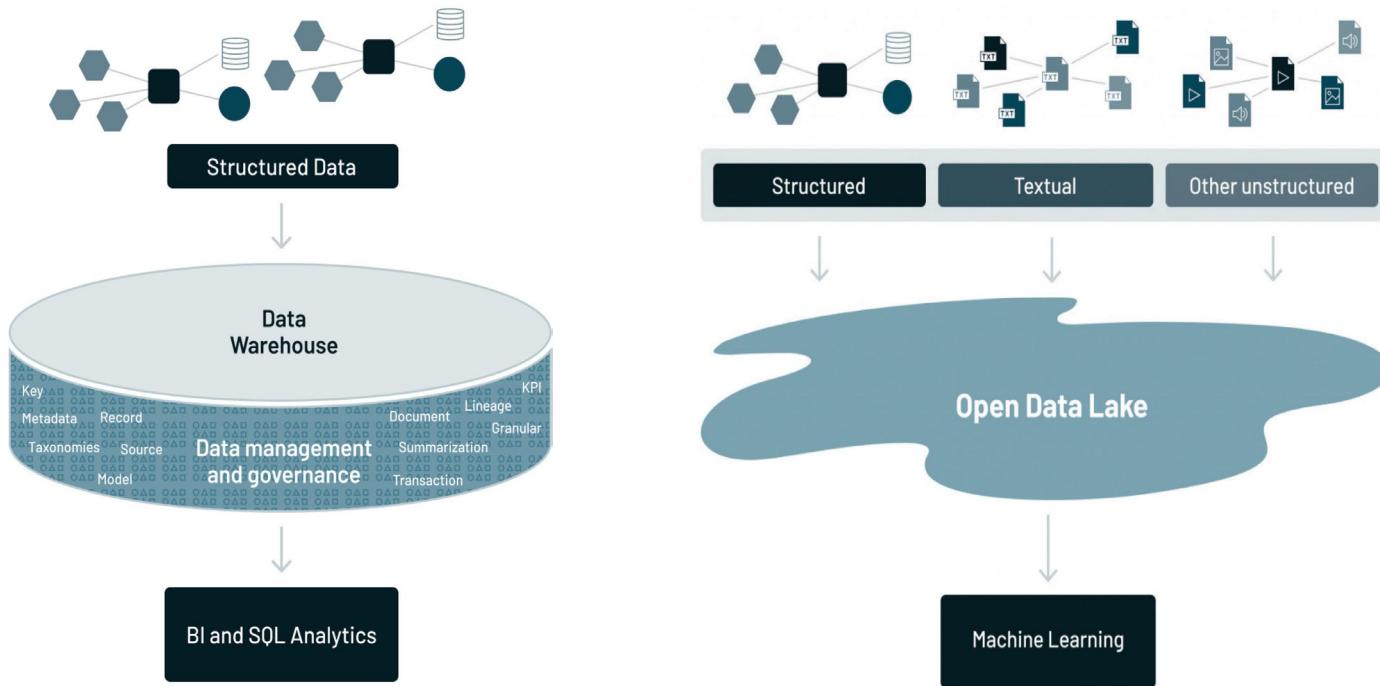
Low utilization, complex per-user config

Databricks Serverless

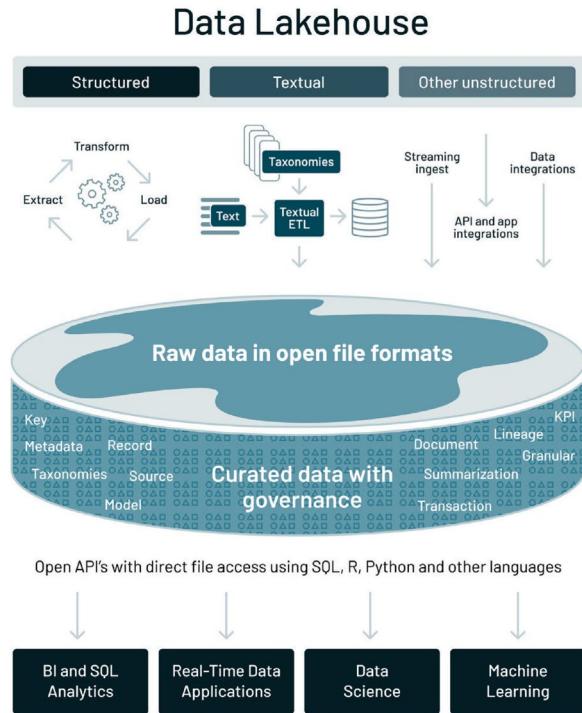


Optimal performance, minimal cost

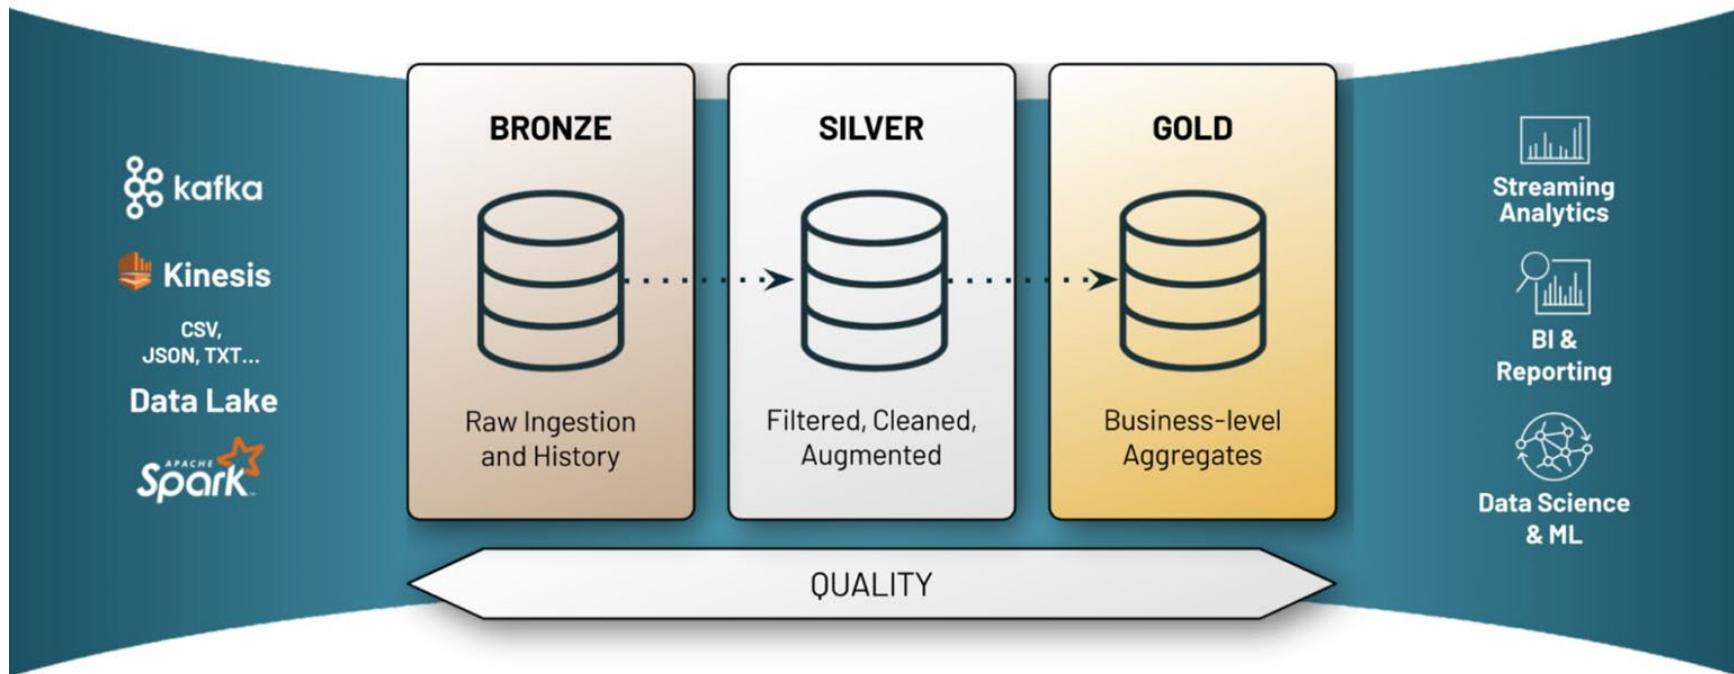
Clash of Architectures



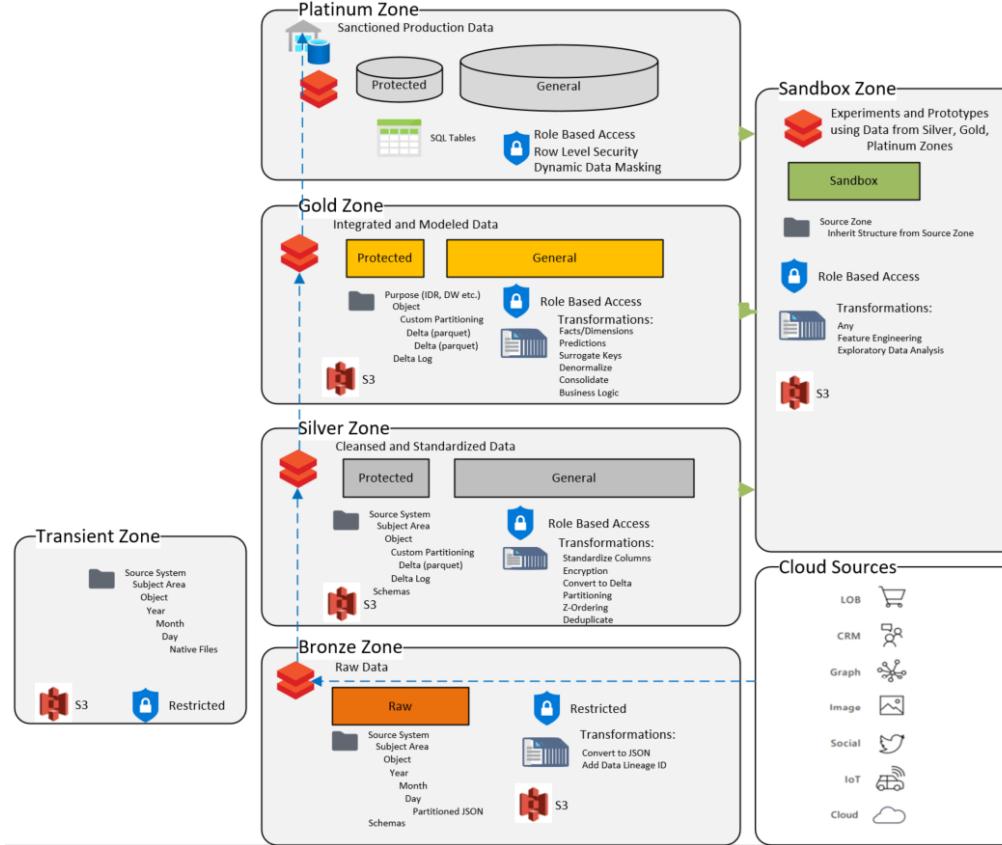
An Evolution of Architectures



Delta Live Tables



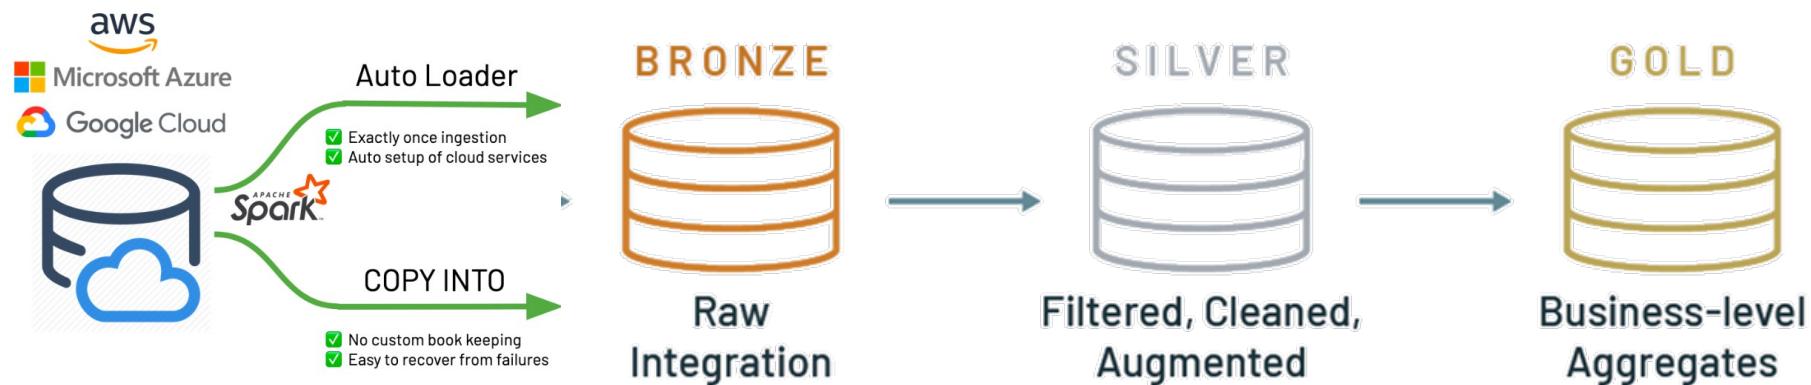
Data Lake Zones



Data Lake Data Sources

Structured Streaming

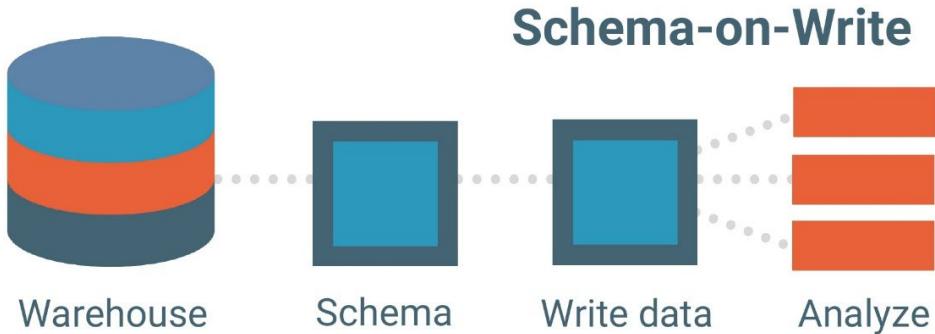
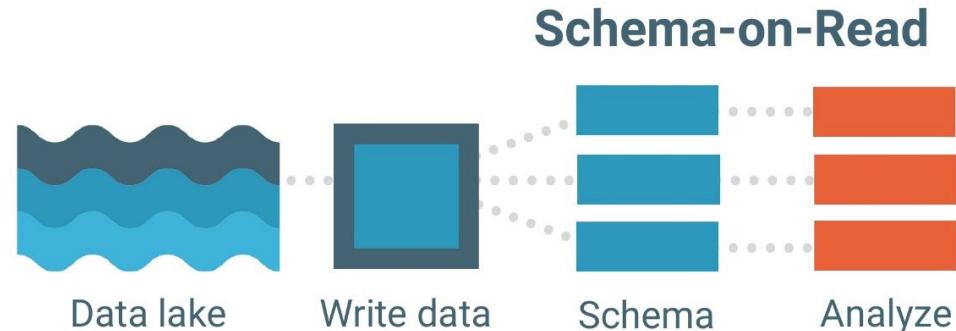
- Transform Avro payload
 - Protocol buffers
 - Apache Kafka
 - AWS Kinesis
 - Google Cloud DataFlow
- Apache Flink
 - Azure Event Hubs
 - Azure Synapse with Structured Streaming
 - Databricks Auto-loader
 - Azure Stream Analytics



Data Lake Enrichment - Schema Management

Common Schema Management Scenarios

- Schema Validation
- Schema Exception Routing
- Schema Evolution
- JSON Schema



Azure Deep Learning Platforms - Databricks



AZURE
DATABRICKS

- Multiple Cloud Support
- Deeper Spark IP
- Cross-cloud support
- Single node support
- Maturer autoscale
- Tighter MLFlow integration
- Native SQL
- Delta Live Tables
- Azure Data Factory integration
- Notebook as Dashboard
- SQL Dashboards/Queries/Alerts
- Unity Catalog
- Cloudfile/Auto-Loader
- Notebook based AutoML

Hive Support
Java/Scala/Python
Spark Scaleout
GPU Support
SparkML Support

- SynapseML support
- Serverless or Dedicated SQL Pool
- Data warehousing
- Data Lake Tables
- Tighter Power BI integration
- .NET language support
- Big Data Scale Model Scoring with Zero Data Movement
- Transact-SQL compatibility
- Azure Purview Integration
- Azure Data Factory workspace integration



AZURE
SYNAPSE

Unity Catalog

Unity Catalog is a unified governance solution for all data and AI assets including files, tables and machine learning models in your lakehouse.



Centralized governance for data and AI

With a common governance model based on ANSI SQL, centrally govern files, tables, dashboards and ML models on any cloud.



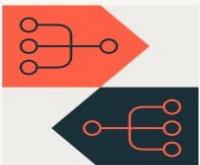
Built-in data search and discovery

Quickly find, understand and reference data from across your data estate, boosting productivity.



Performance and scale

Benefit from enhanced query performance with low-latency metadata serving and auto-tuning of the tables.



Automated lineage for all workloads

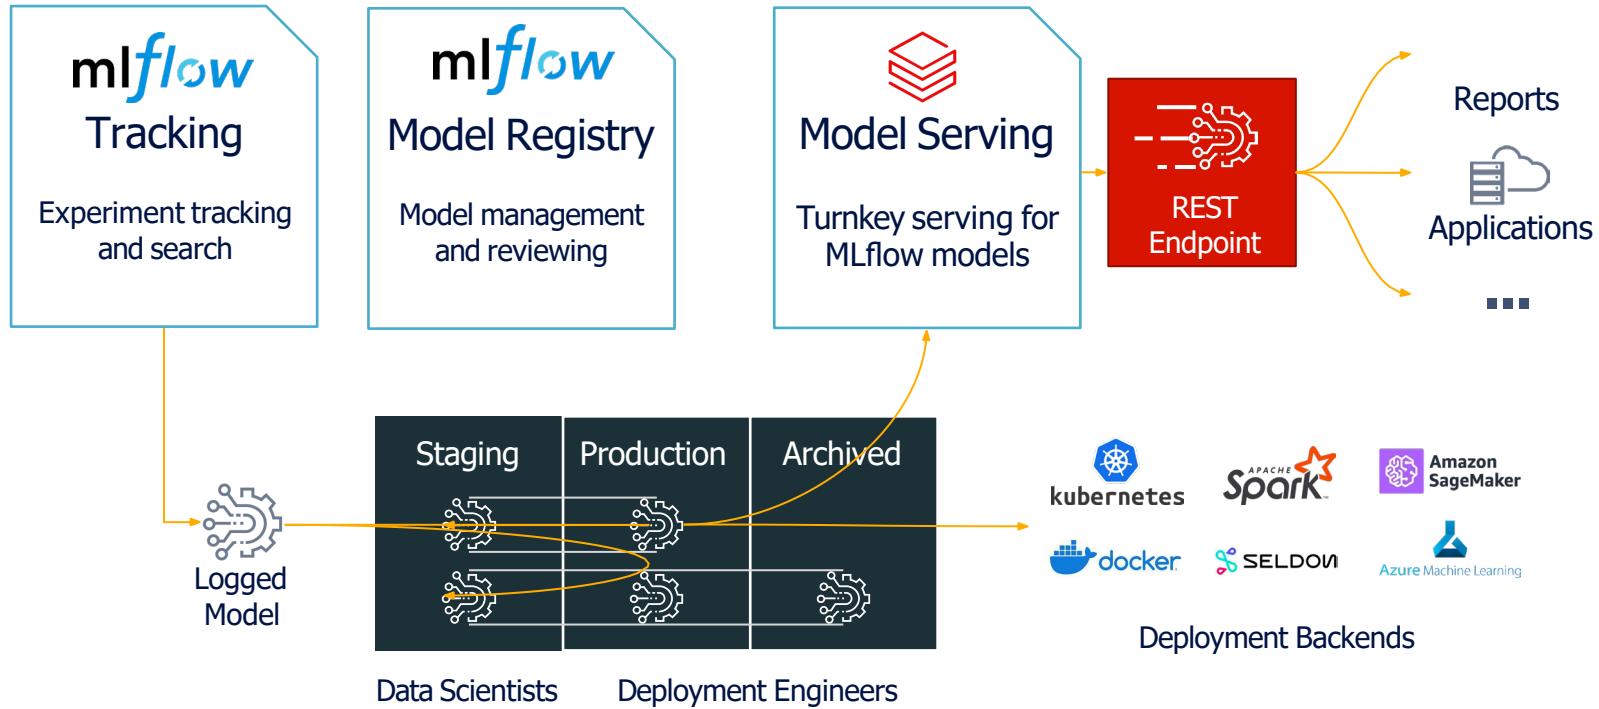
Create a unified, transparent view of your entire data ecosystem with automated and granular lineage for all workloads in SQL, R, Python, Scala and across all asset types — tables, notebooks, workflows and dashboards.



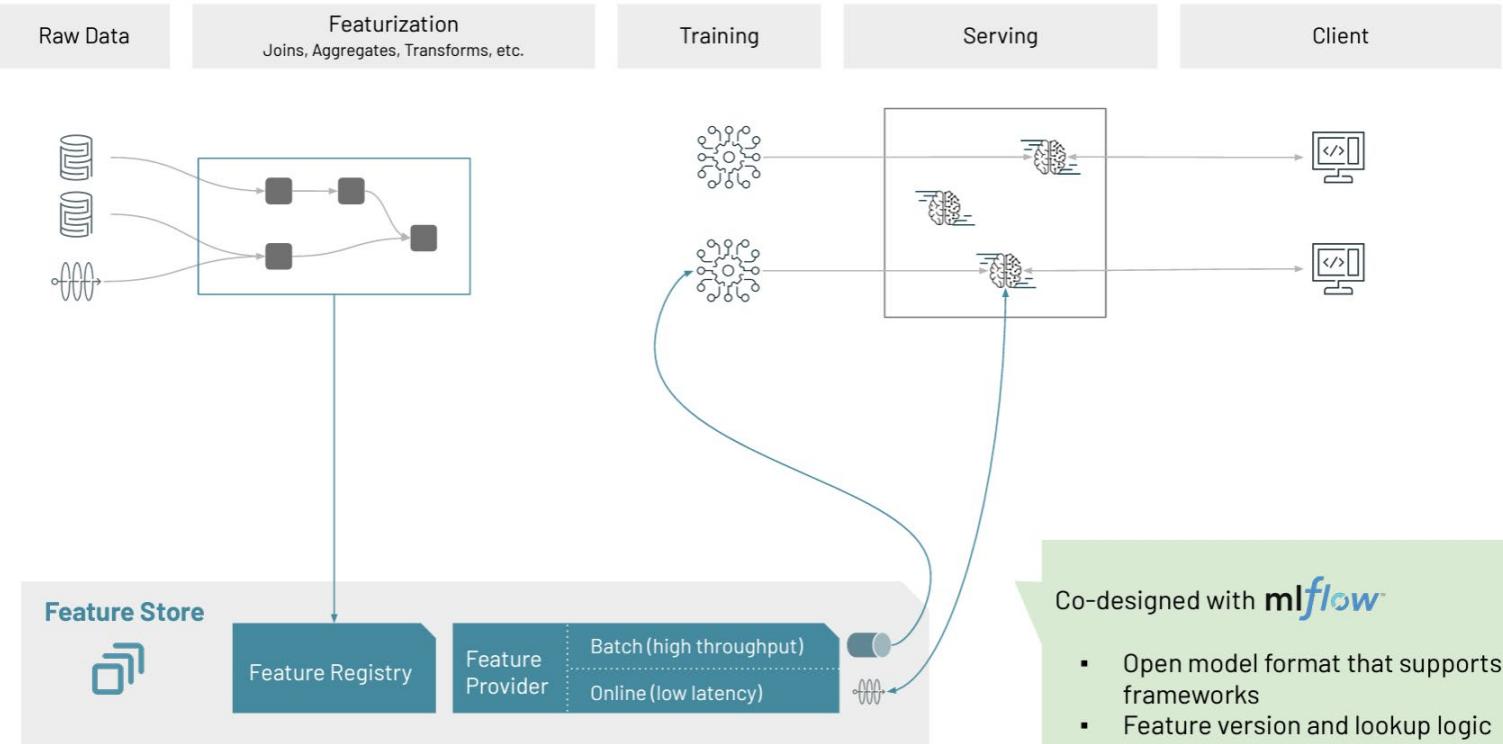
Integrated with your existing tools

Future-proof your data and AI governance with the flexibility to leverage your existing data catalogs and governance solutions.

Open Source MLOps - MLFlow



Data Lake Enrichment - AI/ML Modeling



Co-designed with **mlflow**

- Open model format that supports all ML frameworks
- Feature version and lookup logic hermetically logged with Model



Streaming Analytics Basics

Batch & Streaming - Overview

Batch vs. Streaming



Bucket



Faucet

Batch processing

- High-latency, high delay time
- Large batches of data over time
- Hours to days

Stream processing

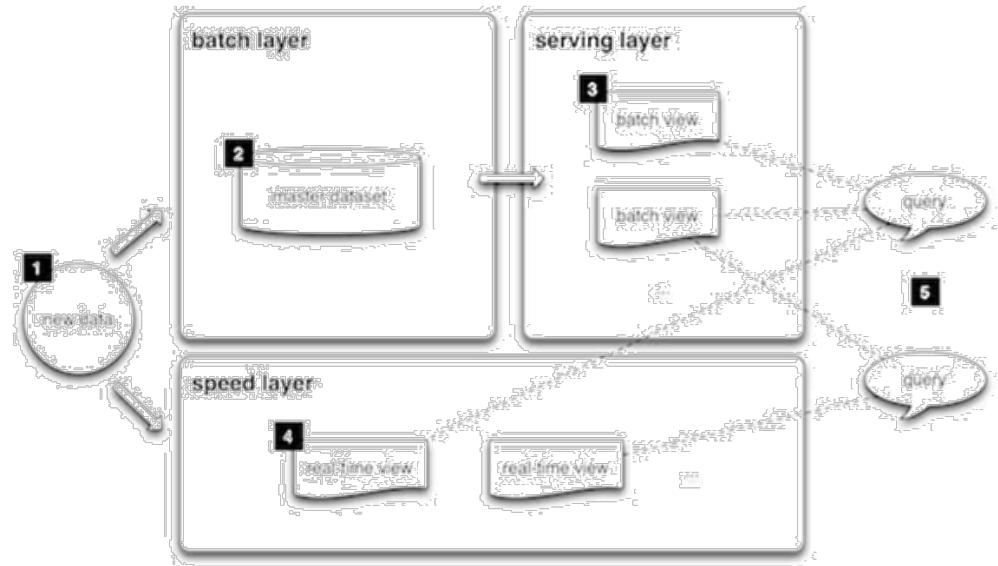
- Low-latency, low delay time
- Micro batches of data at once
- Milliseconds to seconds

<https://www.g2.com/articles/what-is-a-data-lake>

Batch & Streaming - Overview

Lambda Architecture

1. New data arrives to the data lake via batch or streaming
2. Batch layer processes data on disk
3. Serving layer transforms and caches
4. Speed layer processes real time data in memory
5. Clients query for a combination of real time streaming and batch data



Batch & Streaming – Native Execution

Option 1

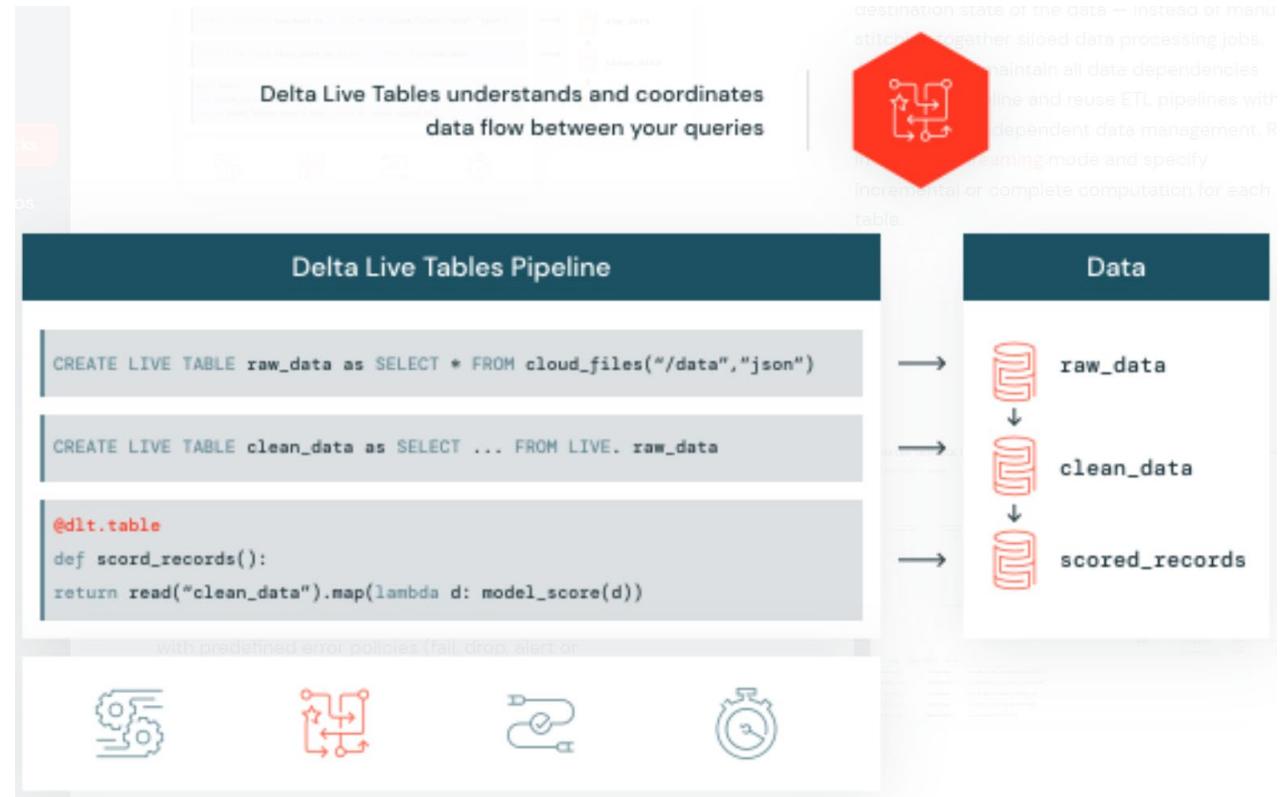
- The combination of using both Batch and Streaming via native Spark gives you:
 - Native performance in unified platform
 - Least amount of storage silos and conversions of data
 - Common skillset of Spark API across Scala, Java, R & Python
 - Lowest cost
 - Simplest set of dependencies



Batch & Streaming – Databricks Managed

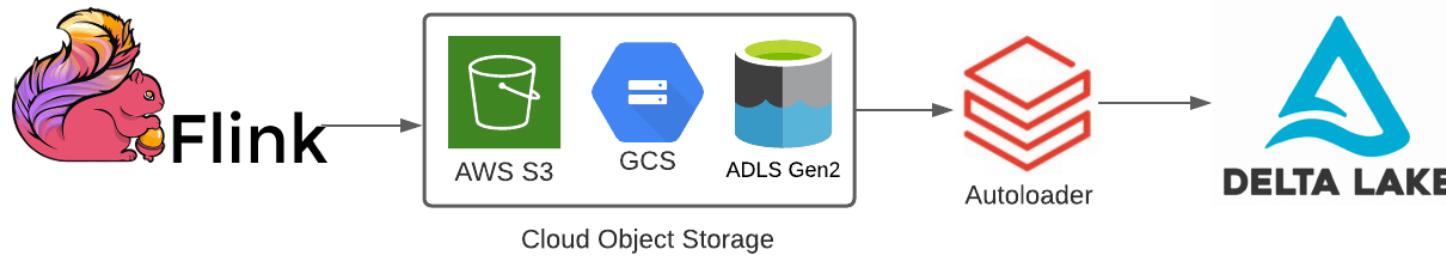
Option 2

- Delta Live Tables
 - Databricks SQL
 - Streaming and
 - Batch seamlessly



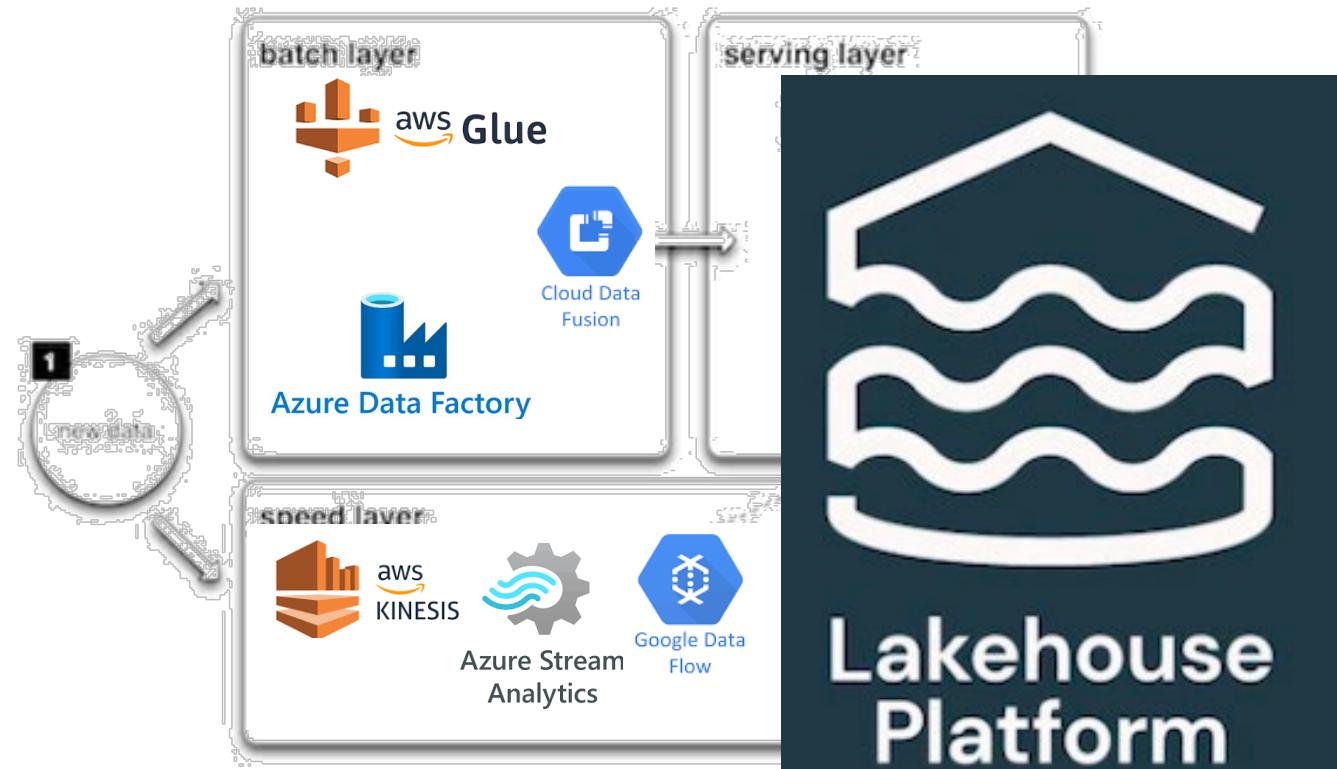
Batch & Streaming - Open Source

Option 3



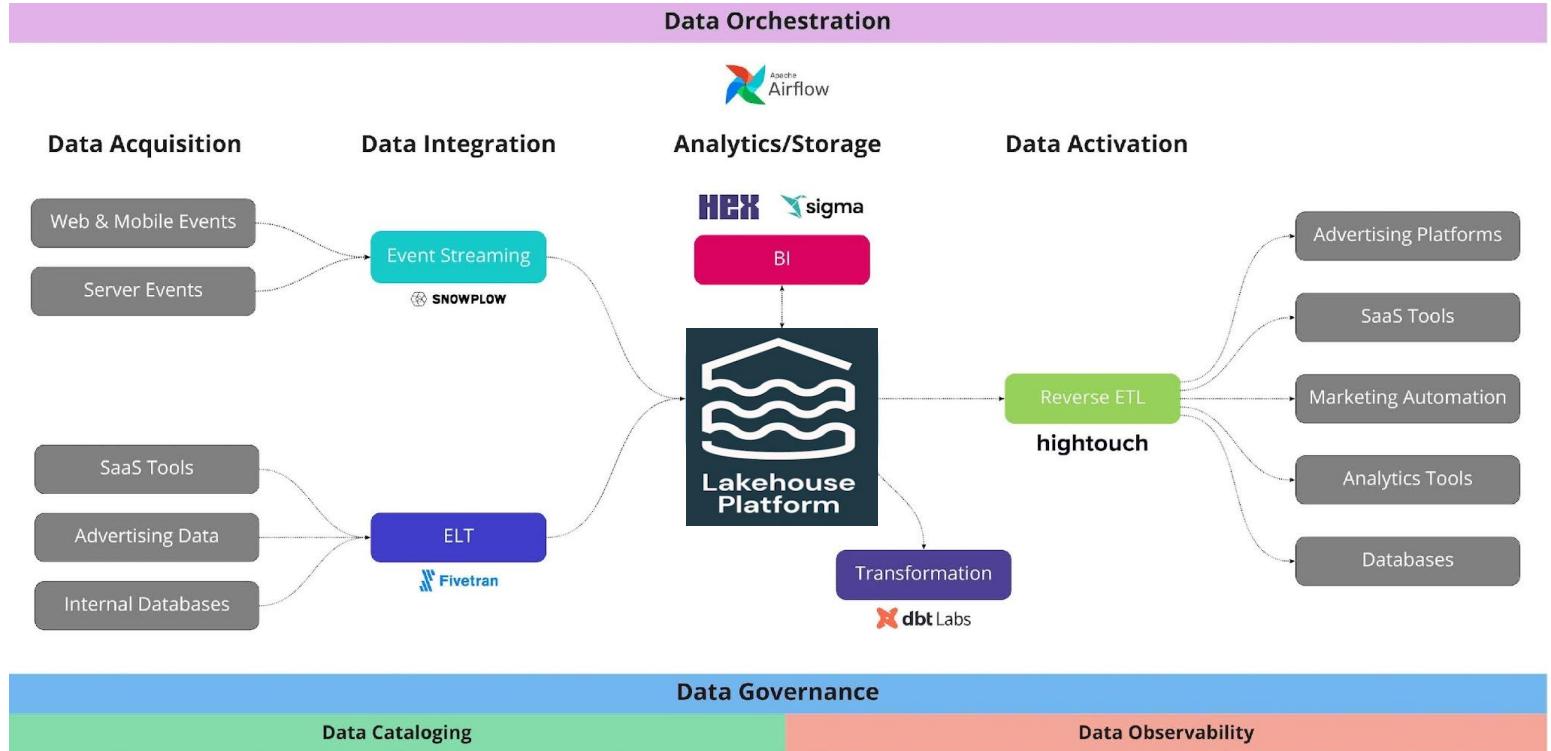
Batch & Streaming - Cloud Native

Option 4



Batch & Streaming – Partner Tools

Option 5



Batch & Streaming – Partner Tools

Data Acquisition

Web & Mobile Events

Server Events

SaaS Tools

Advertising Data

Internal Databases

Data Integration



Analytics/Storage

Data Activation



The process of copying data from your central data warehouse to your operational tools, including but not limited to SaaS tools used for growth, marketing, sales, and support.

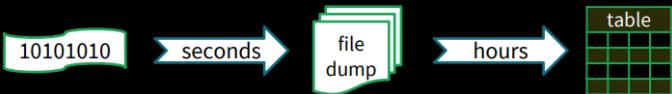
Batch & Streaming – ETL Evolution

Traditional ETL



Raw, dirty, un/semi-structured is data dumped as files

Periodic jobs run every few hours to convert raw data to structured data ready for further analytics



Hours of delay before taking decisions on latest data

Unacceptable when time is of essence
[intrusion detection, anomaly detection, etc.]

Streaming ETL w/ Structured Streaming



Structured Streaming enables raw data to be available as structured data as soon as possible

Complexities in stream processing

Complex Data

Diverse data formats
(json, avro, binary, ...)

Data can be dirty,
late, out-of-order

Complex Workloads

Event time processing

Combining streaming
with interactive queries,
machine learning

Complex Systems

Diverse storage
systems and formats
(SQL, NoSQL, parquet, ...)

System failures

Structured Streaming

stream processing on Spark SQL engine
fast, scalable, fault-tolerant

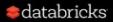
rich, unified, high level APIs
deal with *complex data* and *complex workloads*

rich ecosystem of data sources
integrate with many *storage systems*

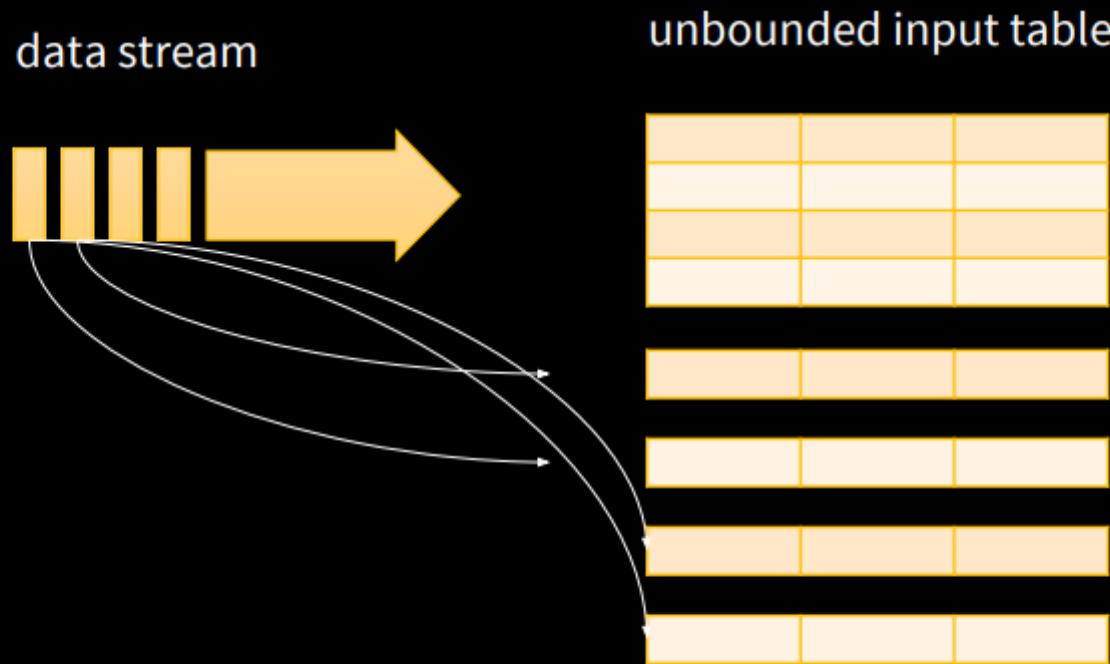
you
should **not** have to
reason about streaming



you
should write simple batch queries
&
Spark
should automatically *streamify*
them



Treat Streams as Unbounded Tables



new data in the
data stream

=

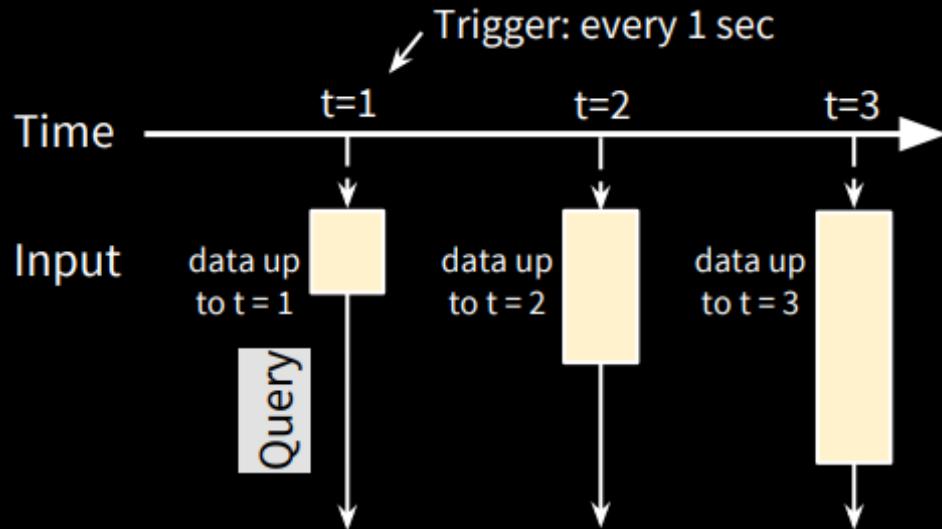
new rows appended
to a unbounded table

New Model

Input: data from source as an append-only table

Trigger: how frequently to check input for new data

Query: operations on input
usual map/filter/reduce
new window, session ops



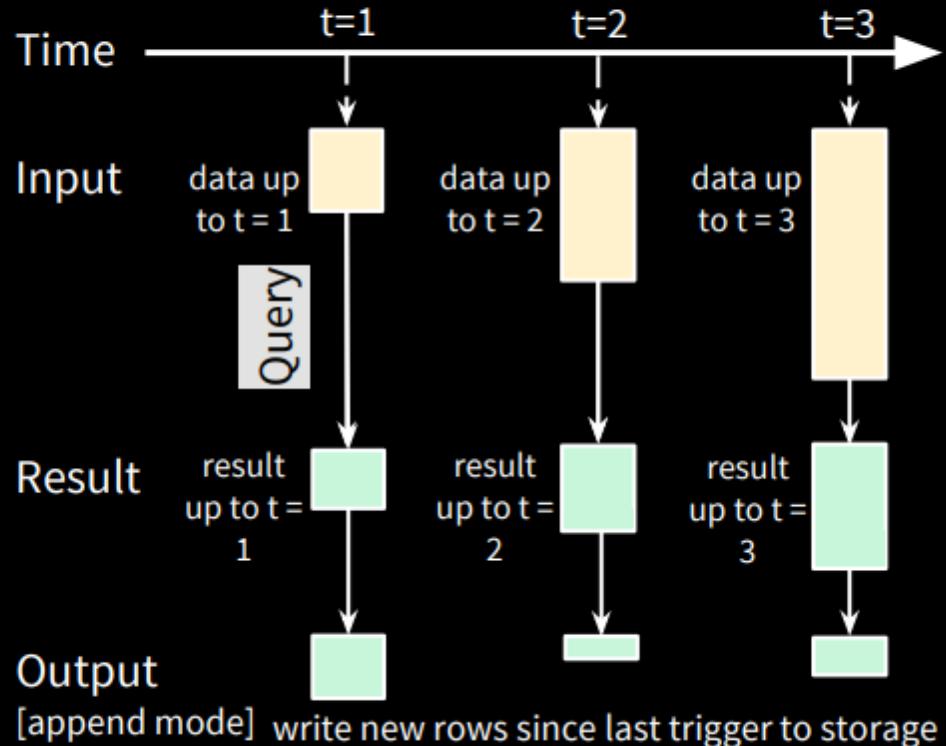
New Model

Result: final operated table
updated after every trigger

Output: what part of result to write
to storage after every trigger

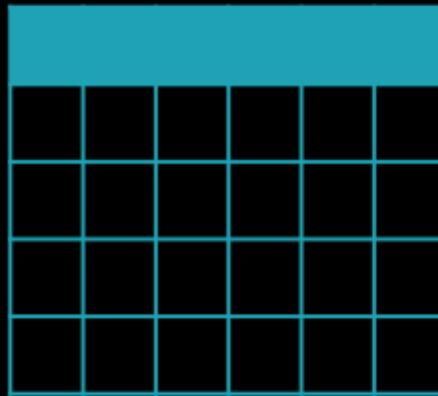
Complete output: write full result table every
time

Append output: write only new rows that got
added to result table since previous batch

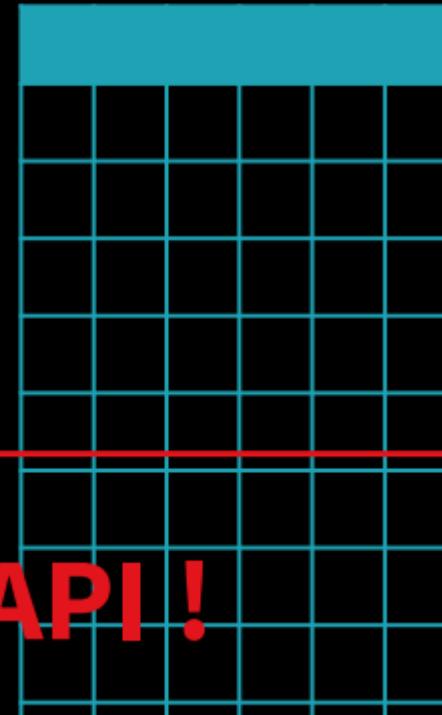


API - *Dataset/DataFrame*

static data =
bounded table



streaming data =
unbounded table



Single API !

Batch Queries with DataFrames

```
input = spark.read  
    .format("json")  
    .load("source-path")  
  
result = input  
    .select("device", "signal")  
    .where("signal > 15")  
  
result.write  
    .format("parquet")  
    .save("dest-path")
```

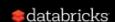
Read from Json file
Select some devices
Write to parquet file



Streaming Queries with DataFrames

```
input = spark.readStream  
    .format("json")  
    .load("source-path")  
  
result = input  
    .select("device", "signal")  
    .where("signal > 15")  
  
result.writeStream  
    .format("parquet")  
    .start("dest-path")
```

Read from Json file stream
Replace `read` with `readStream`
Select some devices
Code does not change
Write to Parquet file stream
Replace `save()` with `start()`

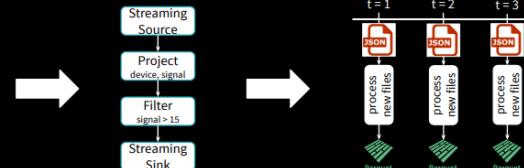


Spark automatically streamifies!

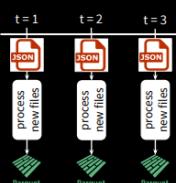
```
input = spark.readStream  
    .format("json")  
    .load("source-path")  
  
result = input  
    .select("device", "signal")  
    .where("signal > 15")  
  
result.writeStream  
    .format("parquet")  
    .start("dest-path")
```

DataFrames,
Datasets,
SQL

Logical Plan



Series of
Incremental



Spark SQL converts batch-like query to a series of
incremental execution plans operating on new batches of
data



Fault-tolerance with Checkpointing

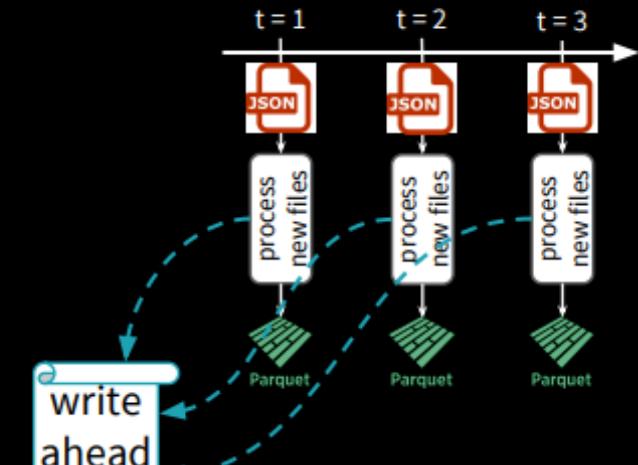
Checkpointing - metadata

(e.g. offsets) of current batch stored
in a *write ahead log* in HDFS/S3

Query can be restarted from the log

Streaming sources can replay the
exact data range in case of failure

Streaming sinks can dedup reprocessed
data when writing, idempotent by design



end-to-end
exactly-once
guarantees

Streaming ETL w/ Structured Streaming

Example

- Json data being received in Kafka
- Parse nested json and flatten it
- Store in structured Parquet table
- Get end-to-end failure guarantees

```
val rawData = spark.readStream
  .format("kafka")
  .option("subscribe", "topic")
  .option("kafka.bootstrap.servers", ...)
  .load()

val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json").as("data"))
  .select("data.*")

val query = parsedData.writeStream
  .option("checkpointLocation", "/checkpoint")
  .partitionBy("date")
  .format("parquet")
```

Reading from Kafka [Spark 2.1]

Support Kafka 0.10.0.1

Specify options to configure

How?

```
kafka.bootstrap.servers => broker1
```

```
val rawData = spark.readStream  
  .format("kafka")  
  .option("kafka.bootstrap.servers", ...)  
  .option("subscribe", "topic")  
  .load()
```

What?

```
subscribe      => topic1,topic2,topic3 // fixed list of topics  
subscribePattern => topic*          // dynamic list of topics  
assign         => {"topicA": [0,1]}    // specific partitions
```

Where?

```
startingOffsets => latest(default) / earliest / {"topicA": {"0": 23, "1": 345} }
```

Reading from Kafka

rawData dataframe has
the following columns

```
val rawData = spark.readStream  
  .format("kafka")  
  .option("subscribe", "topic")  
  .option("kafka.bootstrap.servers",...)  
  .load()
```

key	value	topic	partition	offset	timestamp
[binary]	[binary]	"topicA"	0	345	1486087873
[binary]	[binary]	"topicB"	3	2890	1486086721

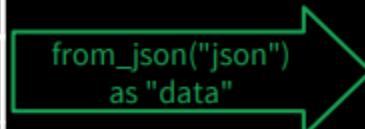
Transforming Data

Cast binary *value* to string
Name it column *json*

```
val parsedData = rawData
    .selectExpr("cast (value as string) as json")
    .select(from_json("json").as("data"))
    .select("data.*")
```

Parse *json* string and expand into
nested columns, name it *data*

json
{ "timestamp": 1486087873, "device": "devA", ...}
{ "timestamp": 1486082418, "device": "devX", ...}



data (nested)		
timestamp	device	...
1486087873	devA	...
1486086721	devX	...

Transforming Data

Cast binary *value* to string
Name it column *json*

```
val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json").as("data"))
  .select("data.*")
```

Parse *json* string and expand
into nested columns, name it
data

Flatten the nested columns

powerful built-in APIs to
perform complex data
transformations

from_json, to_json, explode, ...
100s of functions
(see [our blog post](#))

Writing to Parquet table

Save parsed data as Parquet table in the given path

Partition files by date so that future queries on time slices of data is fast

e.g. query on last 48 hours of data

```
val query = parsedData.writeStream  
    .option("checkpointLocation", ...)  
    .partitionBy("date")  
    .format("parquet")  
    .start("/parquetTable")
```

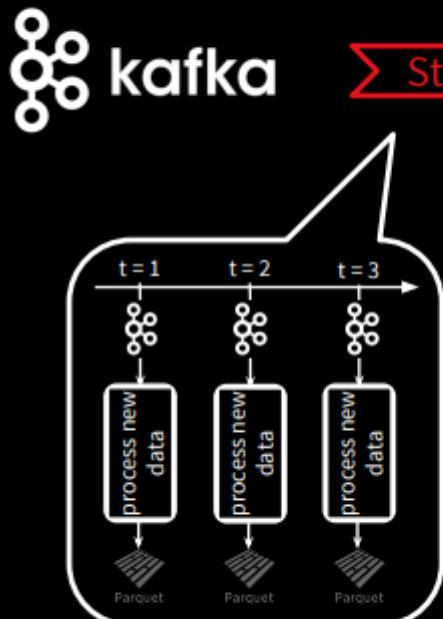
Checkpointing

Enable checkpointing by setting the checkpoint location to save offset logs

`start` actually starts a continuous running `StreamingQuery` in the Spark cluster

```
val query = parsedData.writeStream
  .option("checkpointLocation", ...)
  .format("parquet")
  .partitionBy("date")
  .start("/parquetTable/")
```

Streaming Query

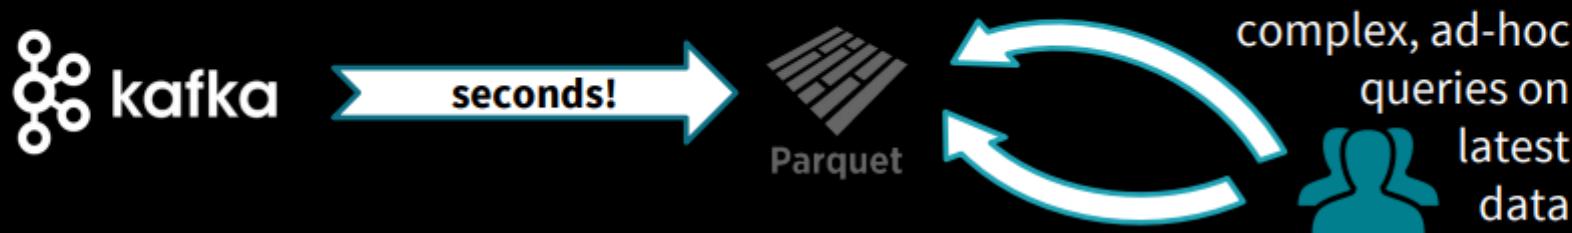


```
val query = parsedData.writeStream  
.option("checkpointLocation", ...)  
.format("parquet")  
.partitionBy("date")  
.start("/parquetTable/")
```

query is a handle to the continuously running StreamingQuery

Used to monitor and manage the execution

Data Consistency on Ad-hoc Queries



Data available for complex, ad-hoc analytics within seconds

Parquet table is updated atomically, ensures *prefix integrity*

Even if distributed, ad-hoc queries will see either all updates from streaming query or none, read more in our blog

<https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html>



Demo, Demo, Demo!!!



THANK YOU

Code & PDF @ <https://github.com/hekaplex/RTDDDMSS>